

# Assignment 1

Niccolò Tosato

3/12/2021

## Section 1: MPI programming

### Section 1.1: 1D ring

The ring is implemented using non-blocking communication routine and 1D virtual topology with periodic boundary. The non-blocking implementation lead a linear growth of the execution time. The execution time of single iteration can be modeled roughly as a *double PingPing*. The real performance will be of course worse than the ideal case, since that the PingPing take into account only 2 processes and does not consider crowded configuration with more than 2 processes. In any case we have a lower bound ideal model to find the expected performance.

#### Software stack and measure setup.

The performance are measured over 10000 iterations, with **UCX** and **InfiniBand**, across *core*, *socket* and *node*. Across two nodes the *mlx5\_0* hardware interface is used. Times for theoretical model come from Intel® MPI benchmark PingPing (Figure 1). Program can be compiled by Makefile with differents options. Timings are taken on rank zero process, it is expected to be slower due topology reason. Follow *ldd* on executable to report specific library used.

```
[s271550@login ring]$ ldd ring.x
linux-vdso.so.1 => (0x00007ffce79ef000)
libm.so.6 => /lib64/libm.so.6 (0x00007f9e9dc88000)
libmpi.so.40 => /opt/area/shared/programs/x86_64/openmpi/4.0.3/gnu/9.3.0/lib/libmpi.so.40 (0x00007f9e9d964000)
```

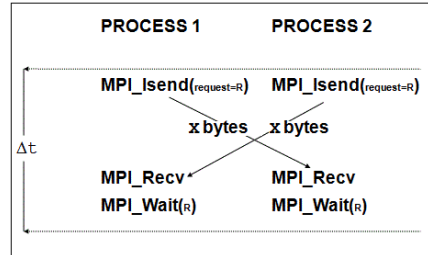


Figure 1: PingPing structure from Intel® MPI Benchmarks User Guide.

#### Map by node model

The network model across two nodes can take into account the latency of the network (dominated by the switch) and the number of processes involved. Each iteration will be lower bounded by the network.

$$Time = N_{procs} \cdot \lambda_{network}$$

The estimated latency between two node with the Intel MPI benchmark is a bit less than the one declared by switch constructor. This lower latency hold only when few communications are going on. With multiple processes involved a more realistic declared latency of 1.35 microseconds lead to a more accurate model, thus the model will be experimentally outperformed when  $N_{procs} = 2$ . When the number of process grow, we can see a slightly slow down of the actual implementation.

Across two node we can estimate optimistically  $\lambda_{network} = 1.01\mu Sec$ .

```
#-----
# Benchmarking PingPing
# #processes = 2
#-----
#
```

#bytes	#repetitions	t[usec]	Mbytes/sec
0	1000	1.00	0.00
1	1000	1.00	1.00
2	1000	1.01	1.98
4	1000	1.01	3.94

We expect for each iteration  $2 \cdot \lambda_{network}$  time, but experimentally this model doesn't hold, the experimental evidence suggest an iteration time of  $\lambda_{network}$ . My guess is that 2 consecutive messages are merged into one unique request routed to the switch. This behaviour can be confirmed using again Intel MPI Benchmark **Exchange**, that implement 1D non periodic ring (Figure 2).

```
#-----
# Benchmarking Exchange
# #processes = 36
#-----
#
```

#bytes	#repetitions	t_min[usec]	t_max[usec]	t_avg[usec]	Mbytes/sec
0	10000	1.46	1.46	1.46	0.00
1	10000	1.45	1.45	1.45	2.75
2	10000	1.45	1.45	1.45	5.50
4	10000	1.46	1.46	1.46	10.97

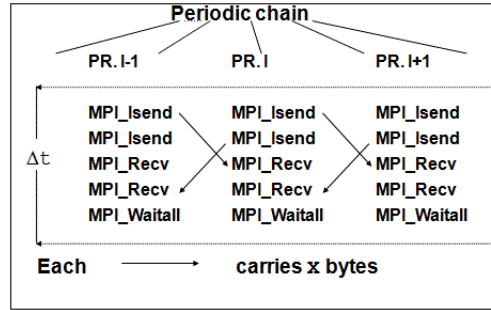


Figure 2: Exchange benchmark structure from Intel® MPI Benchmarks User Guide.

## Map by socket model

With the socket round robin binding selected, we can model the ring behaviour as following, with the usual Intel MPI benchmark estimation.

$$Time = N_{procs} \cdot \lambda_{socket} \cdot 2$$

Again when the sockets become crowded the real performance is worse than the model.

Across two socket we can estimate  $t = 0.49\mu Sec$

```
#-----
# Benchmarking PingPing
# #processes = 2
#-----
#
```

#bytes	#repetitions	t[usec]	Mbytes/sec
0	1000	0.49	0.00
1	1000	0.48	2.06
2	1000	0.49	4.12
4	1000	0.49	8.18

## Map by core model

Mapping by core the processes we have to take into account how many processes are spawned and where. When  $N_{procs}$  is less or equal than the number of core in a single socket, the expected execution time is bounded by the core communication. When the first socket is filled, the first process on the second socket will be the slowest over all process, his neighbors will be placed in the other socket and he will be the slowest in communication. Thus one iteration is long as the slowest process communication time. Experimentally a spike when  $N_{procs} = 13$  is clearly visible. Counter intuitively the performance increase respect  $N_{procs} = 13$  when the number of processes increase and other process are spawned in the second socket.

The slowest processes now have one neighbor on the same socket and one on the other, so the compressive iteration time for the slowest process is  $\lambda_{core} + \lambda_{socket}$ .

$$Time = \begin{cases} N_{procs} \cdot \lambda_{core} \cdot 2 & N_{procs} \leq N_{cpucore} \\ N_{procs} \cdot (\lambda_{core} + \lambda_{socket}) & N_{procs} > N_{cpucore} + 2 \\ N_{procs} \cdot \lambda_{socket} \cdot 2 & N_{procs} = N_{cpucore} + 1 \end{cases}$$

Across two core we can estimate  $t = 0.23\mu Sec$

```
#-----
# Benchmarking PingPing
# #processes = 2
#-----
```

#bytes	#repetitions	t[usec]	Mbytes/sec
0	1000	0.23	0.00
1	1000	0.23	4.27
2	1000	0.24	8.44
4	1000	0.23	17.18

## Experimental results compared with theoretical model

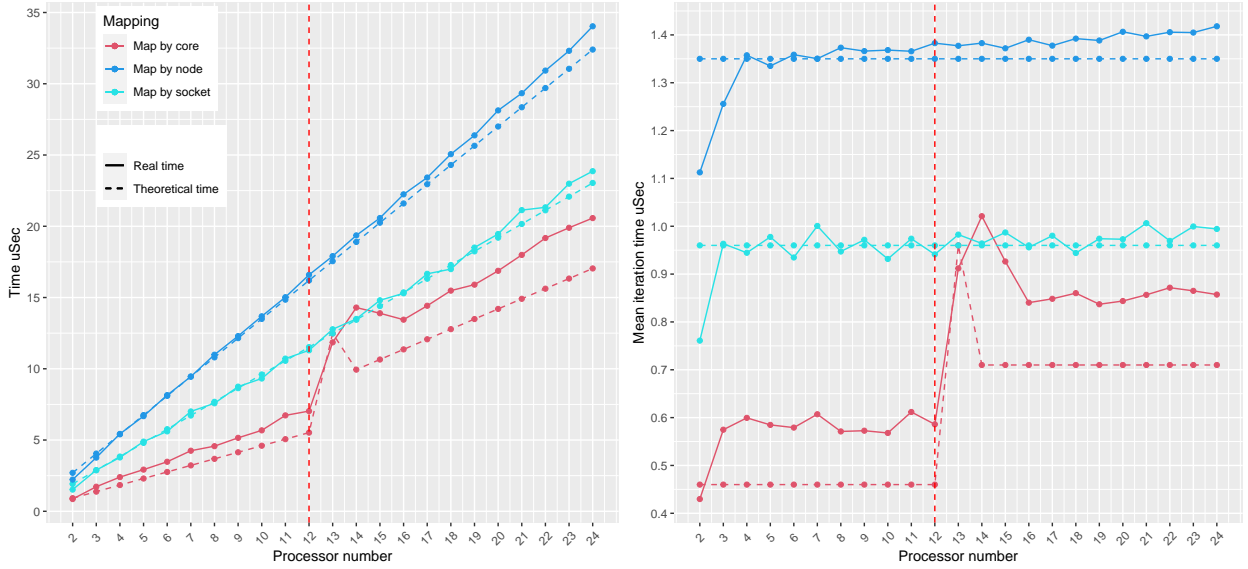


Figure 3: Ring on THIN node up to 24 cores

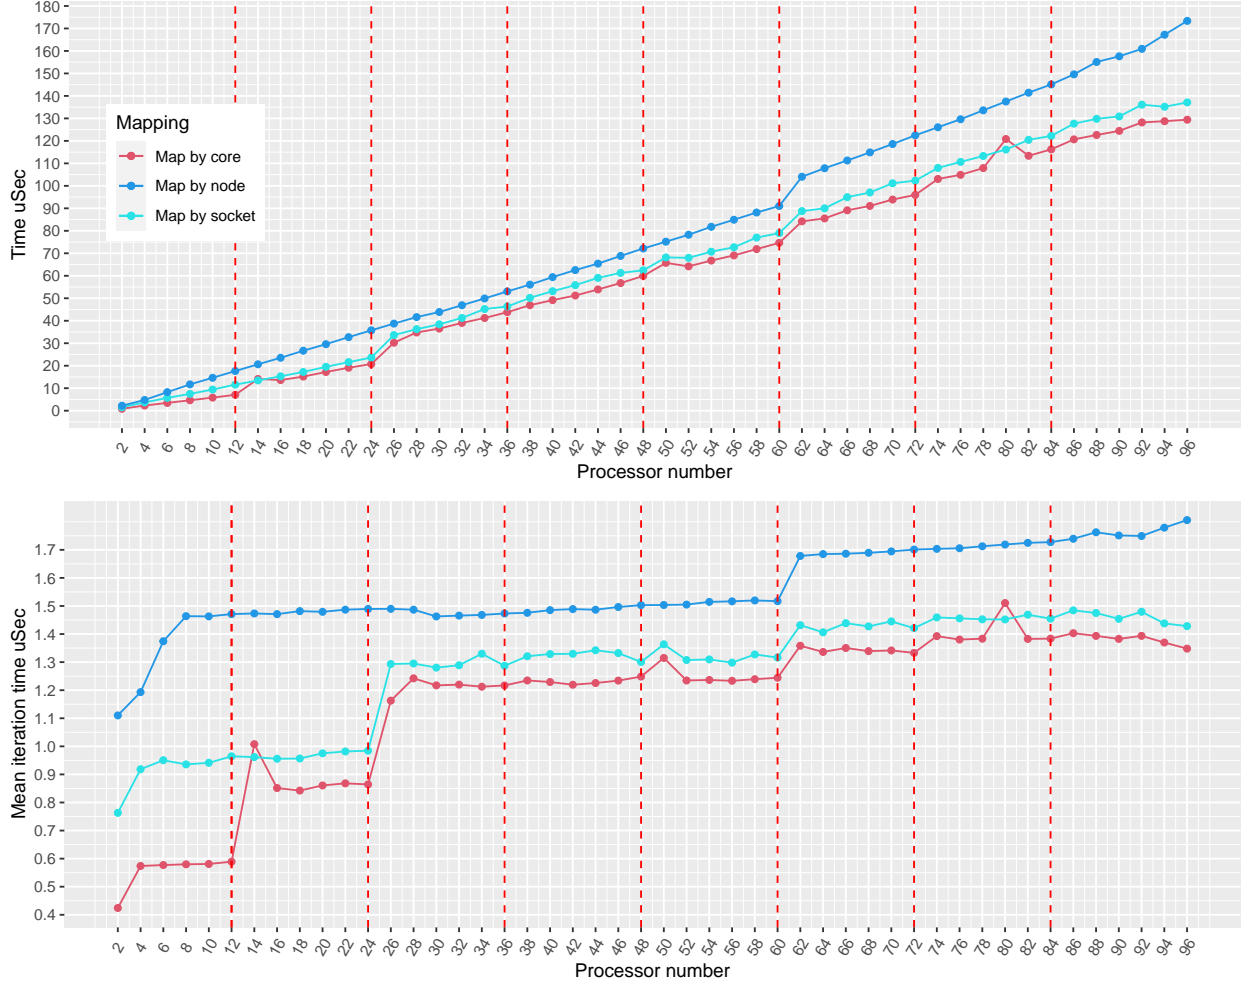


Figure 4: Ring on THIN node up to 96 cores and 4 nodes

Running the ring across 4 node we can recognize the pattern described above. After 24 processes, when socket or node are saturated the changes on mean iteration time are smaller because the communication time between node is the larger and more evident than intra-socket on intra-core, and above 24 processes we have at least one intra-node communication going on, the step confirm that.

## Section 2: measure MPI point to point performance

In order to measure MPI point to point performance, the Intel MPI benchmarks is used. By Intel's documentation it works as follow:




Figure 5: PingPong structure from Intel® MPI Benchmarks User Guide.

The bandwidth and the latency estimate is done across *core*, *socket* and different *node*, combined with different protocols and hardware devices. Each run is repeated 10 times, **openMPI 4.0.3** is used. Entire node was reserved in order to reduce possible source noise in measurement. The **pml** involved in the benchmarks are **OB1** and **UCX**. The **btl** used are **tcp** and **vader**. Across node are selected also different network, with different protocols: 25 Gbit Ethernet and 100 Gbit InfiniBand through Mellanox network switch.

The following graphs show the behaviour inside the same node, i.e mapping the processes across two sockets or in the same socket. Mapping the processes in the same socket show of course often a better performance. The behaviour before the asymptotic plateau is strange and show very different performance among different implementation, the main cause is the cache. Before analyzing the performance we need to know more about the node topology.

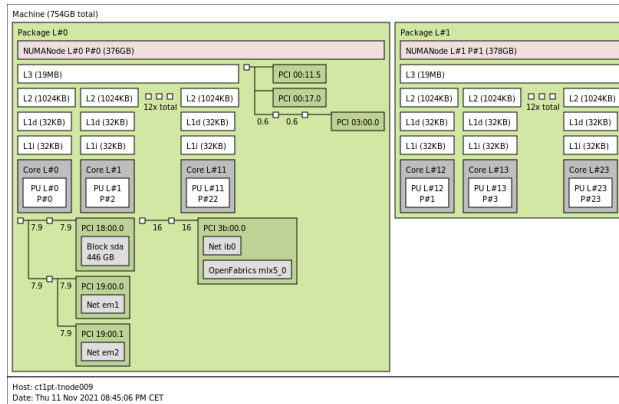






Figure 6: lstopo output on THIN node.


Cache sizes are reported on graph with vertical black dashed line. The behaviour appear strange before 16 MB included, after that size, it is stable because the message size is larger than all caches. The larger cache is **L3** with 19 MB. The **L2** effects, become clear after 1 MB, then all implementations start losing bandwidth. The **L1** effects is not clearly visible from the bandwidth due the latency. To infer with more

accuracy about the cache effect we can perform some profiling test using hardware counters. I modified Intel MPI Benchmark PingPong injecting some code in order to inspect L1 data cache misses and L2 cache misses using PAPI. I tracked all cache misses on MPI\_Send and MPI\_Recv routine of rank 0, after some cycles of warmup. They are reported on absolute value and normalized respect message size. Before 128 KB UCX shows poor performance respect other protocols, after that size there is an huge speed up. This behave will be explained later. Intel InfiniBand also show poor performance respect UCX implementation, this behaviour is explained with cache too.





From first graph, now Intel implementation behaviour is clear, it has the largest cache misses number among all other configurations. The bandwidth spike of **UCX** after 128 KB, can be explained by **L1** and **L2** drop in terms of cache misses. **OB1** implementation seems be the most cache friendly *PML*. Disclaimer: those graphs must be used in qualitative way and not in quantitative way.



This inter-node benchmark reveal the network performance and topology. Two physical network are available, 25 Gbit Ethernet and 100 Gbit InfiniBand. The PCI-E devices are visible from lstopo, we can physically see em1, em2, that are bonded together, creating the interface bond0. To infer about this we can use ifconfig and ip link command.

```
[s271550@ctipt-tnode007 etc]$ ifconfig
bond0: flags=5187<UP,BROADCAST,RUNNING,MASTER,MULTICAST> mtu 1500
        ether 34:80:0d:4e:55:68 txqueuelen 1000 (Ethernet)

br0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
        ether 34:80:0d:4e:55:68 txqueuelen 1000 (Ethernet)

em1: flags=6211<UP,BROADCAST,RUNNING,SLAVE,MULTICAST> mtu 1500
        ether 34:80:0d:4e:55:68 txqueuelen 1000 (Ethernet)

em2: flags=6211<UP,BROADCAST,RUNNING,SLAVE,MULTICAST> mtu 1500
        ether 34:80:0d:4e:55:68 txqueuelen 1000 (Ethernet)

ib0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 2044
Infiniband hardware address can be incorrect! Please read BUGS section in ifconfig(8).
        infiniband 00:00:09:07:FE:80:00: ..... txqueuelen 256 (InfiniBand)


lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
        inet 127.0.0.1 netmask 255.0.0.0
        inet6 ::1 prefixlen 128 scopeid 0x10<host>
```

```
[s271550@ctipt-tnode007 etc]$ ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
2: em1: <BROADCAST,MULTICAST,SLAVE,UP,LOWER_UP> mtu 1500 qdisc mq master bond0 state UP mode DEFAULT group default qlen 1000
3: em2: <BROADCAST,MULTICAST,SLAVE,UP,LOWER_UP> mtu 1500 qdisc mq master bond0 state UP mode DEFAULT group default qlen 1000
4: ib0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 2044 qdisc mq state UP mode DEFAULT group default qlen 256
    link/infiniband ....
5: bond0: <BROADCAST,MULTICAST,MASTER,UP,LOWER_UP> mtu 1500 qdisc noqueue master br0 state UP mode DEFAULT group default qlen 1000
6: br0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default qlen 1000
```



With openMPI implementation and UCX, we can directly select the devices, that lead to a specific protocol as consequence. The devices tested are ib0, br0 and mlx5\_0:1. Ib0 is the IPoIB protocol, br0 lead a TCP communication and mlx5\_0:1 is a pure native InfiniBand device. The theoretical maximum performance are 12.5 GB/s or 12800 MB/s for InfiniBand and 3.125 GB/s or 3200 MB/s for the Ethernet network.

The experimental asymptotic bandwidth measured are:



UCX with InfiniBand and Intel show the best performances among over all configurations in terms of latency and bandwidth. UCX overperform a bit Intel latency. No big cache effects are visible thanks to RDMA that exclude all caches, CPU and kernel action, this explain also the highest possible bandwidth respect core and socket mapping. Real InfiniBand performance is about 95% of theoretical bandwidth, is a very nice result assuming the encoding 64b/66b.

UCX with br0 and OB1 with tcp show comparable latency, but OB1 gain a better bandwidth. The real maximum performance is about 85% of theoretical bandwidth, this is a good result taking into account the heavier tcp protocol respect InfiniBand (ACK,handshake, encoding...). There is then transport overhead and some inefficiency introduced by the cache and CPU (no RDMA available).

IPoIB has a good latency (no Ethernet switch involved) but the bandwidth is not so high as expected. Gpu nodes behave like thin node, no difference are visible, GPU node is a bit slower than thin node, this can be caused by different CPU frequency and node configuration. Cache size are same as thin node, then cache effects are similar. Follow on next page a summary of fitting results on THIN and GPU node.

Table 1: Fitting results

THIN node	Latency uSec	Bandwidth MB/s	GPU node	Latency uSec	Bandwidth MB/s
<b>Core mapping</b>			<b>Core mapping</b>		
UCX IB	0.2440430	5987.179	UCX IB	0.2745335	5652.041
Intel IB	0.2526662	5385.856	Intel IB	0.3247967	5117.350
OB1 vader	0.3658680	3866.456	OB1 vader	0.3850294	3750.822
OB1 tcp	5.0266853	5207.742	OB1 tcp	5.0516070	5150.110
<b>Socket mapping</b>			<b>Socket mapping</b>		
Intel IB	0.5198369	4204.387	Intel IB	0.5685345	4237.763
UCX IB	0.6033689	5123.148	UCX IB	0.6513036	4615.458
OB1 vader	0.8531296	3258.498	OB1 vader	0.8521363	3070.333
OB1 tcp	7.9899043	3325.931	OB1 tcp	8.9058952	3127.998
UCX IB	1.1326157	12182.859			
<b>Node mapping</b>					
Intel ib	1.3189110	12223.147			
UCX ib0	9.6072277	2764.778			
UCX br0	15.8550522	2433.241			
OB1 tcp	16.4297090	2729.666			

## Section 3 : Jacobi solver

### Performance model

In order to predict Jacobi model performance the following model is used:

$$P(L, N) = \frac{L^3 N}{T_s + T_c} [MLUP/s]$$

$L$  represent subdomain size (assuming it as cube), for us it is a costant.  $N$  represent the processes number. The  $L^3 N$  quantity thus is considered the problem size, this amount of work is increased linearly in function of  $N$ , this is a weak scalability. About measure unit: performance is evaluated using  $MLUP/s$ , mega lattice update per sec.

$T_s$  is time swept, is constant and can be estimated using the serial output.

$T_c$  represent the communication time in seconds. This quantity can be modeled simple estimating the latency  $\lambda$ , bandwidth  $B$  and messages size  $C$ .

$$T_c = \frac{C}{B} + 4k\lambda [seconds]$$

About  $C$ , assuming grid point as double:

$$C = L^2 k \cdot 2 \cdot 2 \cdot 8 [byte]$$

About  $k$ : this quantity represent the number of directions in which the halo exchange occurs, it must be multiplied by 2 taking into account the bidirectional halo exchange and again by two taking into account the positive and negative direction.

### Domain decomposition

Several domain decompositions are possible, a priory some of them can be discarded assuming poor performance due buffering, as example between  $(12, 1, 1)$   $(1, 12, 1)$   $(1, 1, 12)$ , the first one is the better and more

memory friendly decomposition, it has the lowest buffering needs. This can be proved experimentally, the first configuration behave better than the other in terms of MLUP/s.

## Experimentals results

### Thin node, core mapping.

$\lambda = 0.24[usec]$  and  $B = 5987[MB/s]$ .

## Usually it is recommended to use `column_spec` before `collapse_rows`, especially in LaTeX, to get a des.

N° procs	Nx	Ny	Nz	k	C MB	Tc s	Theo perf MLUP/s	Real perf MLUP/s	Elapsed time s	NP(1)/P(N)
<b>1</b>	1	1	1	0	0.00000	0.0000000	112.4590	112.2617	54.26	1.000000
<b>4</b>	4	1	1	1	14.95361	0.0024986	449.4678	448.0708	54.44	1.002178
	2	2	1	2	29.90723	0.0049973	449.1002	448.9819	54.37	1.000145
<b>8</b>	8	1	1	1	14.95361	0.0024986	898.9357	894.6112	54.59	1.003893
	4	2	1	2	29.90723	0.0049973	898.2005	896.9956	54.61	1.001224
	2	2	2	3	44.86084	0.0074959	897.4664	894.3671	54.61	1.004167
<b>12</b>	12	1	1	1	14.95361	0.0024986	1348.4035	1329.5301	55.27	1.013246
	4	3	1	2	29.90723	0.0049973	1347.3007	1315.6552	55.70	1.023931
	3	2	2	3	44.86084	0.0074959	1346.1997	1324.6981	55.42	1.016942
	6	2	1	2	29.90723	0.0049973	1347.3007	1323.8481	55.46	1.017595

### Thin node, socket mapping.

$\lambda = 0.6[usec]$  and  $B = 5123[MB/s]$ .

## Usually it is recommended to use `column_spec` before `collapse_rows`, especially in LaTeX, to get a des.

N° procs	Nx	Ny	Nz	k	C MB	Tc s	Theo perf MLUP/s	Real perf MLUP/s	Elapsed time s	NP(1)/P(N)
<b>1</b>	1	1	1	0	0.00000	0.0000000	112.4590	112.3160	54.24	1.000000
<b>4</b>	4	1	1	1	14.95361	0.0029213	449.4056	448.5776	54.35	1.001530
	2	2	1	2	29.90723	0.0058426	448.9760	448.3072	54.36	1.002134
<b>8</b>	8	1	1	1	14.95361	0.0029213	898.8112	895.9560	54.44	1.002871
	4	2	1	2	29.90723	0.0058426	897.9520	895.1857	54.49	1.003734
	2	2	2	3	44.86084	0.0087640	897.0944	895.8977	54.47	1.002936
<b>12</b>	12	1	1	1	14.95361	0.0029213	1348.2169	1332.8755	54.83	1.011191
	4	3	1	2	29.90723	0.0058426	1346.9280	1336.4717	54.84	1.008471
	3	2	2	3	44.86084	0.0087640	1345.6416	1342.5132	54.63	1.003932
	6	2	1	2	29.90723	0.0058426	1346.9280	1342.4993	54.62	1.003943

### Thin node, node mapping.

$\lambda = 1.35[usec]$  and  $B = 12182[MB/s]$ .

## Usually it is recommended to use `column_spec` before `collapse_rows`, especially in LaTeX, to get a des.

N° procs	Nx	Ny	Nz	k	C MB	Tc s	Theo perf MLUP/s	Real perf MLUP/s	Elapsed time s	NP(1)/P(N)
<b>1</b>	1	1	1	0	0.00000	0.0000000	112.459	112.1941	54.26	1.000000
<b>12</b>	12	1	1	1	14.95361	0.0012329	1348.963	1334.3096	54.80	1.009008
	4	3	1	2	29.90723	0.0024658	1348.418	1334.8244	54.75	1.008619
	3	2	2	3	44.86084	0.0036988	1347.874	1342.3393	54.57	1.002973
	6	2	1	2	29.90723	0.0024658	1348.418	1341.9843	54.58	1.003238
<b>24</b>	24	1	1	1	14.95361	0.0012329	2697.926	2680.5702	54.80	1.004510
	12	2	1	2	29.90723	0.0024658	2696.836	2683.9718	54.75	1.003237
	8	3	1	2	29.90723	0.0024658	2696.836	2678.9662	54.78	1.005111
	6	4	1	2	29.90723	0.0024658	2696.836	2676.7929	54.79	1.005927
	6	2	2	3	44.86084	0.0036988	2695.747	2681.9062	54.82	1.004009
	4	3	2	3	44.86084	0.0036988	2695.747	2674.0858	54.85	1.006946
<b>48</b>	48	1	1	1	14.95361	0.0012329	5395.852	5260.9379	56.39	1.023642
	24	2	1	2	29.90723	0.0024658	5393.672	5247.7545	56.51	1.026214
	12	4	1	2	29.90723	0.0024658	5393.672	5249.9587	56.58	1.025783
	12	2	2	3	44.86084	0.0036988	5391.494	5215.2751	56.73	1.032605
	8	6	1	2	29.90723	0.0024658	5393.672	5222.1200	56.67	1.031251
	6	4	2	3	44.86084	0.0036988	5391.494	5244.1732	56.52	1.026915

Even if this model seems accurate, some comments are required: mapping by socket lead a slightly better performance respect core mapping and this is not what expected by theoretical model. If socket map on one hand lead to poor communication performance on other hand exploit better the memory allocation. When one socket is filled with 12 processes the entire grid is stored on socket competence LRDIMM, but spreading processes across 2 socket lead to better exploit of bandwidth using different memory controller and then all memory channel available. This solver at each iteration elaborates an huge amount of memory, given  $L = 1200$  and 12 worker at least 300 GB of ram are elaborated (half memory LRDIMM are saturated, then all single socket channel are used). Also with map-by node option experimentally we obtain better results when several processes are involved. The following results report different domain decomposition and binding with  $L = 1200$ ,  $N = 12$  on thin node. Is reported the maximum value over all runs.

Mapping	Nx	Ny	Nz	MLUP/s
<b>core</b>	12	1	1	1335
	1	1	12	1332
<b>socket</b>	12	1	1	1347
	1	1	12	1345

Performance model doesn't take into account the buffering requirement, given a domain decomposition, needed to exchange the halo point and is not possible theoretically infer about same domain decomposition but on different direction. Experimentally differences on performance are negligible but still present.

**Gpu node, core mapping.**

N° procs	Nx	Ny	Nz	k	C MB	Tc s	Theo perf MLUP/s	Real perf MLUP/s	Elapsed time s	NP(1)/P(N)
<b>1</b>	1	1	1	0	0.00000	0.0000000	79.76744	77.4580	76.52	1.000000
<b>12</b>	12	1	1	1	14.95361	0.0026468	956.62047	850.5267	88.24	1.093200
	4	3	1	2	29.90723	0.0052936	956.03236	849.5713	88.14	1.094429
	3	2	2	3	44.86084	0.0079404	955.44497	849.6338	88.03	1.094348
	6	2	1	2	29.90723	0.0052936	956.03236	850.3519	88.10	1.093424
<b>24</b>	24	1	1	1	14.95361	0.0026468	1913.24094	1690.6353	89.12	1.099936
	12	2	1	2	29.90723	0.0052936	1912.06472	1700.2421	88.86	1.093721
	8	3	1	2	29.90723	0.0052936	1912.06472	1699.0319	88.91	1.094500
	6	4	1	2	29.90723	0.0052936	1912.06472	1700.3159	88.96	1.093674
	6	2	2	3	44.86084	0.0079404	1910.88995	1698.9900	88.89	1.094527
	4	3	2	3	44.86084	0.0079404	1910.88995	1699.5484	88.89	1.094168
<b>48</b>	48	1	1	1	14.95361	0.0026468	3826.48188	2538.1989	114.90	1.465284
	24	2	1	2	29.90723	0.0052936	3824.12944	2534.7265	115.10	1.467291
	12	4	1	2	29.90723	0.0052936	3824.12944	2523.0021	115.40	1.474110
	12	2	2	3	44.86084	0.0079404	3821.77989	2550.3071	114.40	1.458327
	8	6	1	2	29.90723	0.0052936	3824.12944	2549.9056	114.50	1.458557
	6	4	2	3	44.86084	0.0079404	3821.77989	2527.2825	115.30	1.471613

**Gpu node, socket mapping.**

N° procs	Nx	Ny	Nz	k	C MB	Tc s	Theo perf MLUP/s	Real perf MLUP/s	Elapsed time s	NP(1)/P(N)
<b>1</b>	1	1	1	0	0.00000	0.0000000	79.76744	77.4580	76.52	1.000000
<b>12</b>	12	1	1	1	14.95361	0.0032428	956.48797	900.1881	83.73	1.032890
	4	3	1	2	29.90723	0.0064856	955.76773	900.3287	83.15	1.032729
	3	2	2	3	44.86084	0.0097285	955.04857	893.0782	83.52	1.041113
	6	2	1	2	29.90723	0.0064856	955.76773	899.9480	83.18	1.033166
<b>24</b>	24	1	1	1	14.95361	0.0032428	1912.97595	1703.9676	88.97	1.091330
	12	2	1	2	29.90723	0.0064856	1911.53546	1701.9177	89.11	1.092644
	8	3	1	2	29.90723	0.0064856	1911.53546	1700.2848	88.88	1.093694
	6	4	1	2	29.90723	0.0064856	1911.53546	1698.8672	88.94	1.094606
	6	2	2	3	44.86084	0.0097285	1910.09714	1699.6994	88.86	1.094070
	4	3	2	3	44.86084	0.0097285	1910.09714	1699.0902	88.90	1.094463
<b>48</b>	48	1	1	1	14.95361	0.0032428	3825.95189	2528.7299	115.50	1.470771
	24	2	1	2	29.90723	0.0064856	3823.07092	2542.1043	114.50	1.463033
	12	4	1	2	29.90723	0.0064856	3823.07092	2531.6678	115.00	1.469064
	12	2	2	3	44.86084	0.0097285	3820.19428	2546.4954	114.30	1.460510
	8	6	1	2	29.90723	0.0064856	3823.07092	2531.7794	114.60	1.468999
	6	4	2	3	44.86084	0.0097285	3820.19428	2542.6490	114.50	1.462719

On gpu node differences on socket and core mapping are even bigger and enhanced, again we can guess the memory as main cause. Gpu node setup even if has higher frequency LDIMM module, doesn't exploit all possible memory channel. When worker number grow over physical core number the real performance, due hyperthreading, goes dramatically down respect the model. Also execution time show poor performance

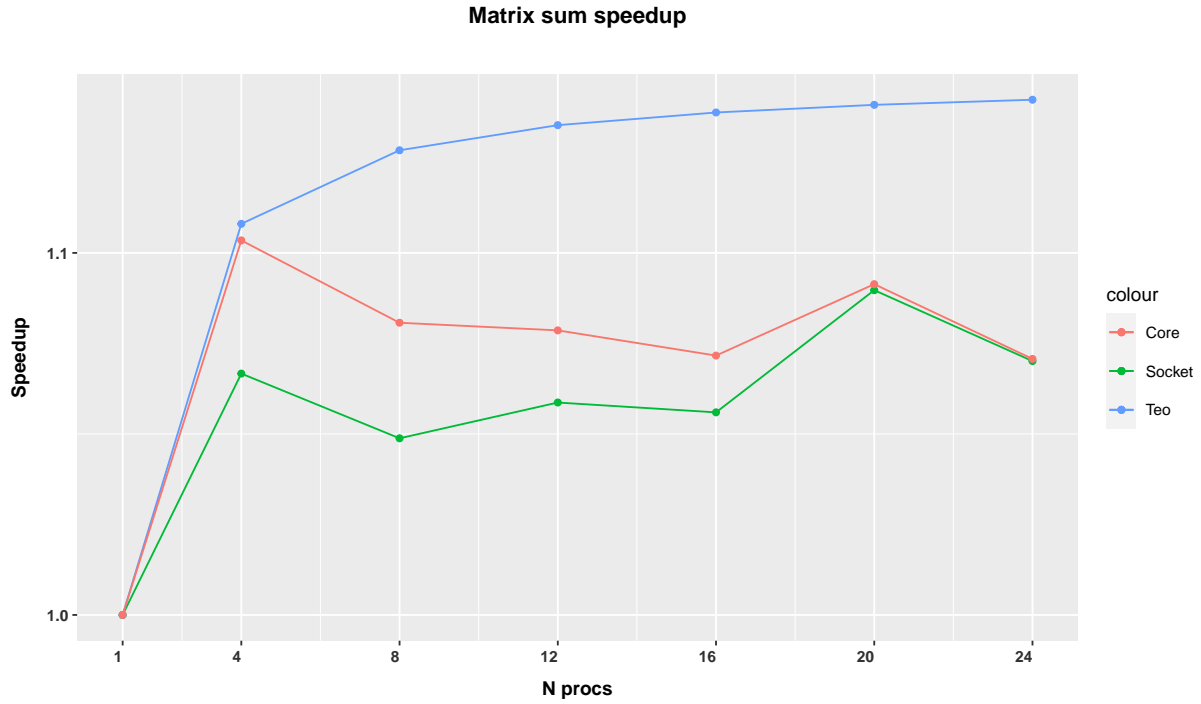
with 48 worker.

Table 2: Fitting results,core and socket map

N° procs	Scatter	Gather	Parallel	Total	N° procs	Scatter	Gather	Parallel	Total
<b>1</b>	16.82058	8.439177	2.635232	49.81	<b>1</b>	16.70198	8.388258	2.620498	49.44
<b>4</b>	13.46379	6.624008	0.657729	45.14	<b>4</b>	13.73519	6.798873	0.638340	46.35
<b>8</b>	14.08911	6.949652	0.335095	46.09	<b>8</b>	14.53400	7.174536	0.329786	47.14
<b>12</b>	14.10777	6.960865	0.223283	46.18	<b>12</b>	14.56708	7.080100	0.218447	46.70
<b>16</b>	15.14609	7.480717	0.161398	46.48	<b>16</b>	15.46470	7.604523	0.162430	46.82
<b>20</b>	14.59633	7.137715	0.131786	45.64	<b>20</b>	14.66246	7.109426	0.127266	45.37
<b>24</b>	15.10795	7.400281	0.109368	46.52	<b>24</b>	15.11786	7.318925	0.108059	46.20

## Matrix sum

Since a 3D array in memory can be represented with a unique linear array, the most effective way to sum it in parallel is using collective operation. The shape of the array doesn't make any difference, also the virtual topology and its relative domain decomposition. We can assume that the problem is not sensible to any topology, since that there is no need of communication between neighbors. The most efficient way is to use MPI\_Scatterv and MPI\_Gatherv collective routine. The expectations for this code are not so high in terms of scalability, matrix sum assuming  $n$  elements each, require  $2n$  memory accesses, a lot of communication and only  $n$  floating point operations. Parallel portion of code then is small respect serial one. Using Ahmdal's law we can make a prediction of speed up. Testing the code with  $2400 \times 1000 \times 700$  matrix using only one processor we can see that the parallel part take only 2.63 seconds over a total runtime of 49 seconds, representing about 5% of execution time (Elapsed is measured using `/usr/bin/time` and detailed Scatter and Gather with `MPI_Wtime()`). Total execution time is comprensive of matrix initialization and error checking. This graph represent then the theoretical maximum speedup supported by Ahmdal's law assuming roughly only 5% of parallel code and communication code and serial code in remaining part and fixed respect processors numbers. This experimentally is a good approximation and catch the trend between [1,24] processes.



I've also implemented a 3D matrix sum program that include a 1D/2D/3D domain decomposition that match virtual topology and as initially requested use collective operation. This program lead to a very high overhead, since that collective operation (in general communication routine) need contiguous memory region to communicate each subdomain to competence worker, to achieve this a unrolling is performed for each subdomain. Once the worker received the subdomain, elaborate it and send back to the master. The master processor need to reordinate the subdomain and assembly the final matrix. Like previous simpler case, this program show a poor performance.



